




expectation,Dynamic Rule Example.Multi-table validation One YAML I would split orders.yaml customer.yaml employee.yaml

- can you write end to end code for by creating yaml and as per the enterprise architecture described above and expectation,Dynamic Rule Example.Multi-table validation

i want to have a dbt style yaml based dQ checks for my hive tables , like null check for a column for a date , date check , count check , etc . for a table with different kinds of partition how can we design a solution which is more robust in nature

can you help me get started with this as per above architecture

I'll help you create an end-to-end DQ framework with YAML-based configurations following your enterprise architecture. This will include dynamic rules, multi-table validation, and partition-aware checks for Hive tables.

Project Structure

text

```
dq_framework/  
├─ config/  
│   ├── tables/  
│   │   ├── orders.yaml  
│   │   ├── customers.yaml  
│   │   └── employees.yaml  
│   └── dq_config.yaml  
└─ src/
```

```
| | └─ parser/  
| |   └─ rule_parser.py  
| | └─ engine/  
| |   └─ metadata_engine.py  
| | └─ resolver/  
| |   └─ partition_resolver.py  
| | └─ compiler/  
| |   └─ rule_compiler.py  
| | └─ generator/  
| |   └─ spark_sql_generator.py  
| | └─ executor/  
| |   └─ spark_executor.py  
| | └─ storage/  
| |   └─ result_store.py  
| └─ notifier/  
|   └─ alert_manager.py  
└─ tests/  
  └─ test_dq_framework.py  
└─ requirements.txt  
└─ main.py
```

1. YAML Configuration Files

config/tables/orders.yaml

```
yaml
```

```
table:  
  name: orders  
  database: sales_db
```

```
description: "Order transactions table"
```

```
schema:
```

```
  columns:
```

- ```
 - name: order_id
 type: string
 nullable: false
```
- ```
    - name: customer_id  
      type: string  
      nullable: false
```
- ```
 - name: order_date
 type: date
 nullable: false
```
- ```
    - name: order_amount  
      type: decimal(10,2)  
      nullable: false
```
- ```
 - name: status
 type: string
 nullable: true
```
- ```
    - name: created_at  
      type: timestamp  
      nullable: false
```

```
partitions:
```

```
  type: date_hierarchy
```

```
  columns:
```

- ```
 - name: year
```

```
 format: yyyy
 - name: month
 format: MM
 - name: day
 format: dd
strategy: dynamic # dynamic/static/full

data_quality:
 rules:
 # Completeness Checks
 - rule_id: ORD_COMP_001
 name: "Order ID Null Check"
 type: null_check
 severity: critical
 columns:
 - order_id
 condition: "order_id IS NULL"
 threshold: 0
 description: "Order ID should never be null"

 - rule_id: ORD_COMP_002
 name: "Required Fields Completeness"
 type: completeness
 severity: critical
 columns:
 - customer_id
 - order_date
 - order_amount
 threshold: 100.0 # 100% completeness expected
 description: "Critical fields must be complete"

 # Freshness Checks
```

```
- rule_id: ORD_FRESH_001
 name: "Order Date Freshness"
 type: freshness
 severity: high
 column: order_date
 max_delay_hours: 24
 description: "Orders should not be older than 24 hours"

Validity Checks
- rule_id: ORD_VAL_001
 name: "Order Amount Positive"
 type: value_range
 severity: high
 column: order_amount
 min_value: 0
 description: "Order amount should be positive"

- rule_id: ORD_VAL_002
 name: "Valid Order Status"
 type: accepted_values
 severity: medium
 column: status
 accepted_values:
 - PENDING
 - PROCESSING
 - COMPLETED
 - CANCELLED
 threshold: 95.0

Uniqueness Checks
- rule_id: ORD_UNIQ_001
 name: "Unique Order ID"
```

```
 type: uniqueness
 severity: critical
 columns:
 - order_id
 threshold: 100.0

Volume Checks
- rule_id: ORD_VOL_001
 name: "Daily Order Volume"
 type: row_count
 severity: warning
 min_count: 100
 max_count: 1000000
 description: "Daily order count should be within range"

Custom SQL Checks
- rule_id: ORD_CUST_001
 name: "Customer-Order Referential Integrity"
 type: custom_sql
 severity: critical
 sql: |
 SELECT COUNT(*) as failed_count
 FROM {database}.orders o
 LEFT JOIN {database}.customers c
 ON o.customer_id = c.customer_id
 WHERE c.customer_id IS NULL
 threshold: 0
 description: "All orders must reference valid customers"

Cross-table Validation
- rule_id: ORD_CROSS_001
 name: "Order-Employee Cross Validation"
```

```
 type: cross_table_validation
 severity: high
 reference_tables:
 - employees
 - departments
 sql: |
 SELECT COUNT(*) as failed_count
 FROM {database}.orders o
 JOIN {database}.employees e ON o.processed_by = e.employee_id
 LEFT JOIN {database}.departments d ON e.dept_id = d.dept_id
 WHERE d.department_status != 'ACTIVE'
 threshold: 0

Partition-specific rules
partition_rules:
 - rule_id: ORD_PART_001
 name: "Partition Completeness"
 type: partition_completeness
 severity: critical
 lookback_days: 7
 expected_partitions: all

 - rule_id: ORD_PART_002
 name: "Partition Size Check"
 type: partition_size
 severity: warning
 min_size_mb: 1
 max_size_mb: 1024
```

## config/tables/customers.yaml



yaml

```
table:
 name: customers
 database: sales_db
 description: "Customer master data"

schema:
 columns:
 - name: customer_id
 type: string
 nullable: false
 - name: customer_name
 type: string
 nullable: false
 - name: email
 type: string
 nullable: false
 - name: signup_date
 type: date
 nullable: false
 - name: customer_segment
 type: string
 nullable: true

partitions:
 type: static
 columns:
 - name: load_date
 format: yyyy-MM-dd

data_quality:
```

```
rules:
 - rule_id: CUST_COMP_001
 name: "Customer Email Format"
 type: regex_match
 severity: high
 column: email
 pattern: "^[a-zA-Z0-9._%+-]+@[a-zA-Z0-9.-]+\.[a-zA-Z]{2,}$"
 threshold: 98.0

 - rule_id: CUST_VAL_001
 name: "Valid Customer Segments"
 type: accepted_values
 severity: medium
 column: customer_segment
 accepted_values:
 - PREMIUM
 - STANDARD
 - BASIC
 - TRIAL
 threshold: 90.0

 - rule_id: CUST_UNIQ_001
 name: "Unique Email Check"
 type: uniqueness
 severity: high
 columns:
 - email
 threshold: 100.0
```

**config/dq\_config.yaml**

```
yaml

framework:
 name: "Enterprise DQ Framework"
 version: "1.0.0"

spark:
 app_name: "DQ_Framework"
 master: "yarn"
 deploy_mode: "cluster"
 executor_memory: "4g"
 executor_cores: 2
 num_executors: 10
 driver_memory: "2g"

hive:
 metastore_uri: "thrift://hive-metastore:9083"
 default_database: "sales_db"

partitions:
 parallel_execution: true
 max_parallel_partitions: 5
 partition_filter:
 start_date: "2024-01-01"
 end_date: "2024-12-31"

storage:
 results:
 type: "hive"
 database: "dq_results"
 table: "dq_check_results"
 partition_column: "check_date"
```

```
metrics:
 type: "postgresql"
 host: "metrics-db.example.com"
 port: 5432
 database: "dq_metrics"
 table: "rule_metrics"

alerting:
 enabled: true
 channels:
 email:
 enabled: true
 smtp_host: "smtp.example.com"
 smtp_port: 587
 recipients:
 critical: ["data-engineering@example.com", "oncall@example.com"]
 high: ["data-engineering@example.com"]
 medium: ["data-analysts@example.com"]

 slack:
 enabled: true
 webhook_url: "https://hooks.slack.com/services/xxx"
 channels:
 critical: "#data-alerts-critical"
 high: "#data-alerts"

 teams:
 enabled: false

pagerduty:
 enabled: true
```

```
 api_key: "pd_api_key"
 service_id: "data-quality"

thresholds:
 failure_tolerance:
 critical: 0 # Zero tolerance for critical failures
 high: 5 # 5% tolerance
 medium: 10
 warning: 20

notification:
 cooldown_minutes: 30
 group_alerts: true
 max_alerts_per_hour: 50
```

## 2. Core Framework Implementation

### src/parser/rule\_parser.py

```
python

import yaml
from pathlib import Path
from typing import Dict, List, Any
from dataclasses import dataclass
from datetime import datetime
import jsonschema
import logging

logger = logging.getLogger(__name__)
```

```
@dataclass
class DQRule:
 rule_id: str
 name: str
 rule_type: str
 severity: str
 table: str
 database: str
 columns: List[str]
 threshold: float
 description: str
 parameters: Dict[str, Any]

class RuleParser:
 """Parses YAML-based DQ rules with validation"""

 def __init__(self, config_path: str):
 self.config_path = Path(config_path)
 self.rules: List[DQRule] = []
 self.tables = {}

 def load_yaml_config(self, file_path: str) -> Dict:
 """Load and validate YAML configuration"""
 try:
 with open(file_path, 'r') as f:
 config = yaml.safe_load(f)
 self._validate_yaml_structure(config)
 return config
 except Exception as e:
 logger.error(f"Failed to load YAML config {file_path}: {e}")
 raise
```

```
def _validate_yaml_structure(self, config: Dict):
 """Validate YAML structure against schema"""
 schema = {
 'type': 'object',
 'required': ['table'],
 'properties': {
 'table': {
 'type': 'object',
 'required': ['name', 'database', 'data_quality'],
 'properties': {
 'name': {'type': 'string'},
 'database': {'type': 'string'},
 'data_quality': {
 'type': 'object',
 'required': ['rules'],
 'properties': {
 'rules': {'type': 'array'}
 }
 }
 }
 }
 }
 }
 try:
 jsonschema.validate(config, schema)
 except jsonschema.exceptions.ValidationError as e:
 raise ValueError(f"Invalid YAML structure: {e}")

def parse_all_configs(self) -> List[DQRule]:
 """Parse all YAML configs in the config directory"""
 config_dir = self.config_path / 'tables'
```

```
yaml_files = list(config_dir.glob('*.yaml')) + list(config_dir.glob('*.yml'))

for yaml_file in yaml_files:
 logger.info(f"Parsing config: {yaml_file}")
 config = self.load_yaml_config(str(yaml_file))
 self.tables[config['table']['name']] = config['table']

 # Parse rules
 for rule_config in config['table']['data_quality']['rules']:
 rule = self._parse_rule(rule_config, config['table'])
 self.rules.append(rule)

return self.rules

def _parse_rule(self, rule_config: Dict, table_config: Dict) -> DQRule:
 """Parse individual rule configuration"""
 return DQRule(
 rule_id=rule_config['rule_id'],
 name=rule_config['name'],
 rule_type=rule_config['type'],
 severity=rule_config['severity'],
 table=table_config['name'],
 database=table_config['database'],
 columns=rule_config.get('columns', []),
 threshold=rule_config.get('threshold', 100.0),
 description=rule_config.get('description', ''),
 parameters={
 'condition': rule_config.get('condition'),
 'max_delay_hours': rule_config.get('max_delay_hours'),
 'min_value': rule_config.get('min_value'),
 'max_value': rule_config.get('max_value'),
 'accepted_values': rule_config.get('accepted_values'),
```



```
 'pattern': rule_config.get('pattern'),
 'min_count': rule_config.get('min_count'),
 'max_count': rule_config.get('max_count'),
 'sql': rule_config.get('sql'),
 'reference_tables': rule_config.get('reference_tables', []),
 'lookback_days': rule_config.get('lookback_days'),
 }
})
```

## src/engine/metadata\_engine.py

python

```
from typing import Dict, List, Any, Optional
from dataclasses import dataclass
from datetime import datetime, timedelta
import pymysql
from pyspark.sql import SparkSession
import logging
```

```
logger = logging.getLogger(__name__)
```

```
@dataclass
```

```
class TableMetadata:
 database: str
 table_name: str
 columns: List[Dict]
 partitions: List[Dict]
 location: str
 file_format: str
 owner: str
```

```
last_updated: datetime
row_count: Optional[int]
size_bytes: Optional[int]

class MetadataEngine:
 """Fetches and manages Hive table metadata"""

 def __init__(self, spark: SparkSession, hive_metastore_uri: str):
 self.spark = spark
 self.hive_metastore_uri = hive_metastore_uri
 self.metadata_cache = {}

 def get_table_metadata(self, database: str, table_name: str) -> TableMetadata:
 """Fetch complete table metadata from Hive metastore"""
 cache_key = f"{database}.{table_name}"

 # Check cache first
 if cache_key in self.metadata_cache:
 return self.metadata_cache[cache_key]

 try:
 # Use Spark catalog to get metadata
 columns_df = self.spark.sql(f"DESCRIBE {database}.{table_name}")
 columns = [row.asDict() for row in columns_df.collect()]

 # Get extended metadata
 extended_df = self.spark.sql(f"DESCRIBE EXTENDED {database}.{table_name}")
 extended_info = {row['col_name']: row['data_type'] for row in extended_df.colle
ct()}

 # Parse partition information
 partitions_df = self.spark.sql(f"SHOW PARTITIONS {database}.{table_name}")
```

```

 partitions = [{'partition': row['partition']} for row in partitions_df.collect()]

 # Get table stats
 stats = self._get_table_stats(database, table_name)

 metadata = TableMetadata(
 database=database,
 table_name=table_name,
 columns=columns,
 partitions=partitions,
 location=extended_info.get('Location', ''),
 file_format=extended_info.get('InputFormat', ''),
 owner=extended_info.get('Owner', ''),
 last_updated=datetime.now(),
 row_count=stats.get('row_count'),
 size_bytes=stats.get('size_bytes')
)

 # Cache the metadata
 self.metadata_cache[cache_key] = metadata
 return metadata

except Exception as e:
 logger.error(f"Failed to fetch metadata for {database}.{table_name}: {e}")
 raise

def _get_table_stats(self, database: str, table_name: str) -> Dict:
 """Get table statistics"""
 try:
 stats_df = self.spark.sql(f"""
 ANALYZE TABLE {database}.{table_name} COMPUTE STATISTICS

```

```
 """)
 return {
 'row_count': None,
 'size_bytes': None
 }
except:
 return {}

def get_column_datatypes(self, database: str, table_name: str) -> Dict[str, str]:
 """Get column name to datatype mapping"""
 metadata = self.get_table_metadata(database, table_name)
 return {col['col_name']: col['data_type'] for col in metadata.columns
 if not col['col_name'].startswith('#')}

def get_partition_columns(self, database: str, table_name: str) -> List[str]:
 """Get partition column names"""
 metadata = self.get_table_metadata(database, table_name)
 partition_cols = []
 for col in metadata.columns:
 if col.get('col_name', '').startswith('# Partition'):
 break
 return partition_cols

def clear_cache(self):
 """Clear metadata cache"""
 self.metadata_cache.clear()
```

## src/resolver/partition\_resolver.py

python

```
from typing import List, Dict, Any, Iterator, Tuple
from datetime import datetime, timedelta
from enum import Enum
import logging

logger = logging.getLogger(__name__)

class PartitionStrategy(Enum):
 FULL = "full"
 INCREMENTAL = "incremental"
 DYNAMIC = "dynamic"
 CUSTOM = "custom"

class PartitionResolver:
 """Resolves partition ranges for DQ checks"""

 def __init__(self, config: Dict):
 self.config = config
 self.default_lookback_days = 7
 self.partition_filter = config.get('partitions', {})

 def resolve_partitions(
 self,
 table_config: Dict,
 rule_config: Dict,
 last_check_timestamp: datetime = None
) -> List[Dict]:
 """
 Resolve which partitions to check based on strategy
 Returns list of partition dictionaries
 """
```

```
partition_config = table_config.get('partitions', {})
strategy = self._determine_strategy(partition_config, rule_config)

if strategy == PartitionStrategy.FULL:
 return self._resolve_full_partitions(table_config)
elif strategy == PartitionStrategy.INCREMENTAL:
 return self._resolve_incremental_partitions(table_config, last_check_timestamp)
elif strategy == PartitionStrategy.DYNAMIC:
 return self._resolve_dynamic_partitions(table_config, rule_config)
else:
 return self._resolve_custom_partitions(table_config, rule_config)

def _determine_strategy(self, partition_config: Dict, rule_config: Dict) -> PartitionStrategy:
 """Determine partition strategy based on config and rule type"""
 rule_type = rule_config.get('type', '')

 # Different rules might need different strategies
 if rule_type in ['null_check', 'completeness', 'uniqueness']:
 return PartitionStrategy.FULL
 elif rule_type in ['freshness', 'row_count']:
 return PartitionStrategy.DYNAMIC
 elif rule_type == 'partition_completeness':
 return PartitionStrategy.FULL
 else:
 return PartitionStrategy.INCREMENTAL

def _resolve_full_partitions(self, table_config: Dict) -> List[Dict]:
 """Get all partitions"""
 # This would query Hive metastore for all partitions
 # Simplified version returning date range
 start_date = datetime.strptime(
```

```
 self.partition_filter.get('start_date', '2024-01-01'),
 '%Y-%m-%d'
)
 end_date = datetime.strptime(
 self.partition_filter.get('end_date', datetime.now().strftime('%Y-%m-%d')),
 '%Y-%m-%d'
)

 partitions = []
 current = start_date
 while current <= end_date:
 partitions.append({
 'year': current.strftime('%Y'),
 'month': current.strftime('%m'),
 'day': current.strftime('%d')
 })
 current += timedelta(days=1)

 return partitions

def _resolve_incremental_partitions(
 self,
 table_config: Dict,
 last_check_timestamp: datetime = None
) -> List[Dict]:
 """Get partitions since last check"""
 if last_check_timestamp is None:
 last_check_timestamp = datetime.now() - timedelta(hours=24)

 end_date = datetime.now()

 partitions = []
```

```
current = last_check_timestamp
while current <= end_date:
 partitions.append({
 'year': current.strftime('%Y'),
 'month': current.strftime('%m'),
 'day': current.strftime('%d')
 })
 current += timedelta(days=1)

return partitions

def _resolve_dynamic_partitions(
 self,
 table_config: Dict,
 rule_config: Dict
) -> List[Dict]:
 """Resolve partitions based on rule requirements"""
 lookback_days = rule_config.get('lookback_days', self.default_lookback_days)
 end_date = datetime.now()
 start_date = end_date - timedelta(days=lookback_days)

 partitions = []
 current = start_date
 while current <= end_date:
 partitions.append({
 'year': current.strftime('%Y'),
 'month': current.strftime('%m'),
 'day': current.strftime('%d')
 })
 current += timedelta(days=1)

 return partitions
```



```

def _resolve_custom_partitions(
 self,
 table_config: Dict,
 rule_config: Dict
) -> List[Dict]:
 """Custom partition resolution logic"""
 # Implementation depends on specific requirements
 return self._resolve_dynamic_partitions(table_config, rule_config)

def generate_partition_filters(self, partitions: List[Dict]) -> List[str]:
 """Generate SQL partition filter clauses"""
 filters = []
 for partition in partitions:
 conditions = []
 for key, value in partition.items():
 conditions.append(f"{key}='{value}'")
 filters.append(" AND ".join(conditions))
 return filters

```

## src/compiler/rule\_compiler.py

python

```

from typing import Dict, List, Any
from dataclasses import dataclass
import logging

logger = logging.getLogger(__name__)

@dataclass

```

```
class CompiledRule:
 rule_id: str
 name: str
 rule_type: str
 severity: str
 sql_query: str
 expected_threshold: float
 partition_info: Dict

class RuleCompiler:
 """Compiles YAML rules into executable SQL queries"""

 def __init__(self):
 self.rule_templates = {
 'null_check': self._compile_null_check,
 'completeness': self._compile_completeness,
 'freshness': self._compile_freshness,
 'value_range': self._compile_value_range,
 'accepted_values': self._compile_accepted_values,
 'uniqueness': self._compile_uniqueness,
 'row_count': self._compile_row_count,
 'regex_match': self._compile_regex_match,
 'custom_sql': self._compile_custom_sql,
 'cross_table_validation': self._compile_cross_table_validation,
 }

 def compile_rule(
 self,
 rule_config: Dict,
 table_config: Dict,
 partition_filter: str = None
) -> CompiledRule:
```

```
"""Compile a single rule into executable SQL"""
rule_type = rule_config['type']

if rule_type not in self.rule_templates:
 raise ValueError(f"Unknown rule type: {rule_type}")

compiler_func = self.rule_templates[rule_type]
sql_query = compiler_func(rule_config, table_config, partition_filter)

return CompiledRule(
 rule_id=rule_config['rule_id'],
 name=rule_config['name'],
 rule_type=rule_type,
 severity=rule_config['severity'],
 sql_query=sql_query,
 expected_threshold=rule_config.get('threshold', 0),
 partition_info={'filter': partition_filter} if partition_filter else {}
)

def _add_partition_clause(self, query: str, partition_filter: str) -> str:
 """Add partition filter to query if provided"""
 if partition_filter:
 if 'WHERE' in query.upper():
 return query.replace('WHERE', f'WHERE {partition_filter} AND')
 elif 'from' in query.lower():
 # Find the table reference and add WHERE after it
 return query + f' WHERE {partition_filter}'
 return query

def _compile_null_check(
 self, rule: Dict, table: Dict, partition_filter: str = None
) -> str:
```

```

 """Compile null check rule"""
 columns = rule['columns']
 database = table['database']
 table_name = table['name']

 null_checks = " OR ".join([f"{col} IS NULL" for col in columns])

 query = f"""
 SELECT
 '{rule['rule_id']}' as rule_id,
 COUNT(*) as total_rows,
 SUM(CASE WHEN {null_checks} THEN 1 ELSE 0 END) as failed_rows,
 ROUND(SUM(CASE WHEN {null_checks} THEN 1 ELSE 0 END) * 100.0 / COUNT(*), 2) as
failure_percentage
 FROM {database}.{table_name}
 """

 return self._add_partition_clause(query, partition_filter)

 def _compile_completeness(
 self, rule: Dict, table: Dict, partition_filter: str = None
) -> str:
 """Compile completeness check for multiple columns"""
 columns = rule['columns']
 database = table['database']
 table_name = table['name']

 completeness_checks = []
 for col in columns:
 completeness_checks.append(f"""
 SUM(CASE WHEN {col} IS NOT NULL THEN 1 ELSE 0 END) as {col}_complete,
 COUNT({col}) as {col}_total
 """)

```

```

 """
 """

 query = f"""
 SELECT
 '{rule['rule_id']}' as rule_id,
 COUNT(*) as total_rows,
 {'',''.join(completeness_checks)}
 FROM {database}.{table_name}
 """

 return self._add_partition_clause(query, partition_filter)

def _compile_freshness(
 self, rule: Dict, table: Dict, partition_filter: str = None
) -> str:
 """Compile freshness check"""
 column = rule['column']
 max_delay_hours = rule.get('max_delay_hours', 24)
 database = table['database']
 table_name = table['name']

 query = f"""
 SELECT
 '{rule['rule_id']}' as rule_id,
 MAX({column}) as max_value,
 MIN({column}) as min_value,
 COUNT(*) as total_rows,
 SUM(CASE
 WHEN {column} < DATE_SUB(CURRENT_TIMESTAMP(), {max_delay_hours} * 3600)
 THEN 1 ELSE 0
 END) as stale_rows
 FROM {database}.{table_name}
 """

```

```

 """

 return self._add_partition_clause(query, partition_filter)

def _compile_value_range(
 self, rule: Dict, table: Dict, partition_filter: str = None
) -> str:
 """Compile value range check"""
 column = rule['column']
 min_value = rule.get('min_value')
 max_value = rule.get('max_value')
 database = table['database']
 table_name = table['name']

 conditions = []
 if min_value is not None:
 conditions.append(f"{column} < {min_value}")
 if max_value is not None:
 conditions.append(f"{column} > {max_value}")

 condition_str = " OR ".join(conditions)

 query = f"""
 SELECT
 '{rule['rule_id']}' as rule_id,
 COUNT(*) as total_rows,
 SUM(CASE WHEN {condition_str} THEN 1 ELSE 0 END) as failed_rows,
 MIN({column}) as min_value,
 MAX({column}) as max_value,
 AVG({column}) as avg_value
 FROM {database}.{table_name}
 WHERE {column} IS NOT NULL
 """

```

```

 """

 return self._add_partition_clause(query, partition_filter)

def _compile_accepted_values(
 self, rule: Dict, table: Dict, partition_filter: str = None
) -> str:
 """Compile accepted values check"""
 column = rule['column']
 accepted_values = rule['accepted_values']
 database = table['database']
 table_name = table['name']

 values_str = ", ".join([f"'{v}'" for v in accepted_values])

 query = f"""
 SELECT
 '{rule['rule_id']}' as rule_id,
 COUNT(*) as total_rows,
 COUNT(DISTINCT {column}) as distinct_values,
 SUM(CASE WHEN {column} NOT IN ({values_str}) THEN 1 ELSE 0 END) as invalid_valu
es,
 COLLECT_SET(CASE WHEN {column} NOT IN ({values_str}) THEN {column} END) as inva
lid_value_list
 FROM {database}.{table_name}
 """

 return self._add_partition_clause(query, partition_filter)

def _compile_uniqueness(
 self, rule: Dict, table: Dict, partition_filter: str = None
) -> str:

```

```
"""Compile uniqueness check"""
columns = rule['columns']
database = table['database']
table_name = table['name']

columns_str = ", ".join(columns)

query = f"""
SELECT
 '{rule['rule_id']}' as rule_id,
 COUNT(*) as total_rows,
 COUNT(DISTINCT {columns_str}) as unique_values,
 COUNT(*) - COUNT(DISTINCT {columns_str}) as duplicate_count
FROM {database}.{table_name}
"""

return self._add_partition_clause(query, partition_filter)

def _compile_row_count(
 self, rule: Dict, table: Dict, partition_filter: str = None
) -> str:
 """Compile row count check"""
 min_count = rule.get('min_count', 0)
 max_count = rule.get('max_count')
 database = table['database']
 table_name = table['name']

 query = f"""
SELECT
 '{rule['rule_id']}' as rule_id,
 COUNT(*) as row_count,
 CASE
```



```

 WHEN COUNT(*) < {min_count} THEN 'BELOW_MIN'
 WHEN COUNT(*) > {max_count} THEN 'ABOVE_MAX'
 ELSE 'OK'
 END as status
FROM {database}.{table_name}
"""

 return self._add_partition_clause(query, partition_filter)

def _compile_regex_match(
 self, rule: Dict, table: Dict, partition_filter: str = None
) -> str:
 """Compile regex pattern match check"""
 column = rule['column']
 pattern = rule['pattern']
 database = table['database']
 table_name = table['name']

 query = f"""
 SELECT
 '{rule['rule_id']}' as rule_id,
 COUNT(*) as total_rows,
 SUM(CASE WHEN {column} IS NOT NULL AND {column} NOT RLIKE '{pattern}'
 THEN 1 ELSE 0 END) as invalid_format_count,
 ROUND(SUM(CASE WHEN {column} IS NOT NULL AND {column} NOT RLIKE '{pattern}'
 THEN 1 ELSE 0 END) * 100.0 / NULLIF(COUNT(*), 0), 2) as failure_percentage
 FROM {database}.{table_name}
 """

 return self._add_partition_clause(query, partition_filter)

def _compile_custom_sql(

```

```

 self, rule: Dict, table: Dict, partition_filter: str = None
) -> str:
 """Compile custom SQL rule"""
 sql = rule['sql']
 database = table['database']

 # Replace placeholders
 sql = sql.replace('{database}', database)

 return self._add_partition_clause(sql, partition_filter)

 def _compile_cross_table_validation(
 self, rule: Dict, table: Dict, partition_filter: str = None
) -> str:
 """Compile cross-table validation rule"""
 sql = rule['sql']
 database = table['database']

 # Replace placeholders for all reference tables
 sql = sql.replace('{database}', database)

 return self._add_partition_clause(sql, partition_filter)

```

## src/generator/spark\_sql\_generator.py

python

```

from typing import Dict, List
from pyspark.sql import SparkSession, DataFrame
import logging

```

```
logger = logging.getLogger(__name__)

class SparkSQLGenerator:
 """Generates and optimizes Spark SQL for DQ checks"""

 def __init__(self, spark: SparkSession):
 self.spark = spark
 self.query_cache = {}

 def generate_check_query(self, compiled_rule: 'CompiledRule') -> str:
 """Generate optimized Spark SQL query for a compiled rule"""
 base_query = compiled_rule.sql_query

 # Add query hints for optimization
 optimized_query = self._add_optimization_hints(base_query)

 # Add result formatting
 final_query = self._format_result_query(optimized_query, compiled_rule)

 return final_query

 def _add_optimization_hints(self, query: str) -> str:
 """Add Spark SQL optimization hints"""
 # Add broadcast hints for small tables
 if 'JOIN' in query.upper():
 query = query.replace('JOIN', '/*+ BROADCAST */ JOIN', 1)

 # Add partition hints
 if 'WHERE' in query.upper() and 'year=' in query.lower():
 query = f"-- Enable partition pruning\n{query}"

 return query
```

```
def _format_result_query(self, query: str, compiled_rule: 'CompiledRule') -> str:
 """Format query to include metadata in results"""
 metadata_fields = f"""
 '{compiled_rule.rule_id}' as rule_id,
 '{compiled_rule.name}' as rule_name,
 '{compiled_rule.rule_type}' as rule_type,
 '{compiled_rule.severity}' as severity,
 current_timestamp() as check_timestamp
 """

 # Check if query already has rule_id
 if 'rule_id' not in query:
 query = f"""
 SELECT
 {metadata_fields},
 result.*
 FROM ({query}) result
 """

 return query

def generate_batch_query(
 self,
 compiled_rules: List['CompiledRule'],
 union_results: bool = True
) -> str:
 """Generate batch query for multiple rules"""
 if not union_results:
 return [self.generate_check_query(rule) for rule in compiled_rules]

 individual_queries = [self.generate_check_query(rule) for rule in compiled_rules]
```

```
 return " UNION ALL ".join(individual_queries)

 def cache_query_result(self, rule_id: str, df: DataFrame):
 """Cache query results for reuse"""
 self.query_cache[rule_id] = df

 def get_cached_result(self, rule_id: str) -> DataFrame:
 """Get cached query result"""
 return self.query_cache.get(rule_id)

 def clear_cache(self):
 """Clear query cache"""
 self.query_cache.clear()
```

## src/executor/spark\_executor.py

python

```
from pyspark.sql import SparkSession, DataFrame
from typing import List, Dict, Any
from datetime import datetime
import logging
from concurrent.futures import ThreadPoolExecutor, as_completed

logger = logging.getLogger(__name__)

class SparkExecutor:
 """Executes DQ checks on Spark with parallel processing"""

 def __init__(self, spark: SparkSession, config: Dict):
```

```
self.spark = spark
self.config = config
self.spark_config = config.get('spark', {})
self._configure_spark()

def _configure_spark(self):
 """Configure Spark session for DQ checks"""
 spark_conf = self.spark.sparkContext.getConf()

 for key, value in self.spark_config.items():
 spark_conf.set(f"spark.{key}", str(value))

 # Enable adaptive query execution
 spark_conf.set("spark.sql.adaptive.enabled", "true")
 spark_conf.set("spark.sql.adaptive.coalescePartitions.enabled", "true")

def execute_single_check(self, query: str) -> DataFrame:
 """Execute a single DQ check"""
 try:
 logger.info(f"Executing query: {query[:200]}...")
 df = self.spark.sql(query)
 df.cache() # Cache for potential reuse
 return df
 except Exception as e:
 logger.error(f"Failed to execute query: {e}")
 raise

def execute_batch_checks(self, queries: List[str], parallel: bool = True) -> List[DataF
rame]:
 """Execute multiple DQ checks"""
 results = []
```

```

if parallel and len(queries) > 1:
 with ThreadPoolExecutor(max_workers=min(len(queries), 5)) as executor:
 future_to_query = {
 executor.submit(self.execute_single_check, query): query
 for query in queries
 }

 for future in as_completed(future_to_query):
 query = future_to_query[future]
 try:
 result = future.result()
 results.append(result)
 except Exception as e:
 logger.error(f"Query failed: {query[:100]}... Error: {e}")
 else:
 for query in queries:
 result = self.execute_single_check(query)
 results.append(result)

return results

def execute_with_retry(
 self,
 query: str,
 max_retries: int = 3,
 retry_delay: int = 60
) -> DataFrame:
 """Execute query with retry logic"""
 for attempt in range(max_retries):
 try:
 return self.execute_single_check(query)
 except Exception as e:

```

```

 if attempt < max_retries - 1:
 logger.warning(f"Query failed (attempt {attempt + 1}/{max_retries}), re
trying... Error: {e}")
 import time
 time.sleep(retry_delay)
 else:
 raise

def get_execution_metrics(self) -> Dict:
 """Get Spark execution metrics"""
 sc = self.spark.sparkContext
 return {
 'active_executors': sc._jsc.sc().getExecutorMemoryStatus().size(),
 'total_tasks': sum(sc.statusTracker().getJobIdsForGroup(None)),
 'executor_memory': self.spark_config.get('executor_memory', '4g'),
 'executor_cores': self.spark_config.get('executor_cores', 2)
 }

def optimize_partition_strategy(self, table_name: str, partition_columns: List[str]):
 """Optimize partition strategy for queries"""
 for col in partition_columns:
 self.spark.sql(f"ALTER TABLE {table_name} RECOVER PARTITIONS")

```

## src/storage/result\_store.py

python

```

from pyspark.sql import SparkSession, DataFrame
from typing import Dict, List
from datetime import datetime
import logging

```



```
logger = logging.getLogger(__name__)

class ResultStore:
 """Stores DQ check results in Hive and PostgreSQL"""

 def __init__(self, spark: SparkSession, config: Dict):
 self.spark = spark
 self.config = config
 self.hive_config = config.get('storage', {}).get('results', {})
 self.pg_config = config.get('storage', {}).get('metrics', {})

 def store_results(self, results_df: DataFrame, check_date: str = None):
 """Store DQ check results in Hive"""
 if check_date is None:
 check_date = datetime.now().strftime('%Y-%m-%d')

 database = self.hive_config.get('database', 'dq_results')
 table = self.hive_config.get('table', 'dq_check_results')

 try:
 # Add metadata columns
 enriched_df = results_df.withColumn('check_date', results_df.check_timestamp.cast('date'))

 # Write to Hive table with partition
 enriched_df.write \
 .mode('append') \
 .partitionBy('check_date') \
 .format('parquet') \
 .saveAsTable(f"{database}.{table}")
```

```
 logger.info(f"Stored {enriched_df.count()} results in {database}.{table}")

 except Exception as e:
 logger.error(f"Failed to store results in Hive: {e}")
 raise

def store_metrics(self, metrics: List[Dict]):
 """Store aggregated metrics in PostgreSQL"""
 try:
 import psycopg2
 from psycopg2.extras import execute_values

 conn = psycopg2.connect(
 host=self.pg_config['host'],
 port=self.pg_config['port'],
 database=self.pg_config['database'],
 user=self.pg_config.get('user', 'dq_user'),
 password=self.pg_config.get('password')
)

 cursor = conn.cursor()

 # Insert metrics
 insert_sql = """
 INSERT INTO dq_metrics
 (rule_id, table_name, check_timestamp, total_rows, failed_rows,
 failure_percentage, partition_info, status)
 VALUES %s
 """

 values = [
 (
```

```
 m['rule_id'],
 m['table_name'],
 m['check_timestamp'],
 m['total_rows'],
 m['failed_rows'],
 m['failure_percentage'],
 m.get('partition_info', ''),
 m['status']
)
 for m in metrics
]

execute_values(cursor, insert_sql, values)
conn.commit()

logger.info(f"Stored {len(metrics)} metrics in PostgreSQL")

except Exception as e:
 logger.error(f"Failed to store metrics in PostgreSQL: {e}")
 if conn:
 conn.rollback()
finally:
 if cursor:
 cursor.close()
 if conn:
 conn.close()

def get_historical_results(
 self,
 rule_id: str,
 days: int = 7
) -> DataFrame:
```

```
"""Get historical results for trend analysis"""
database = self.hive_config.get('database', 'dq_results')
table = self.hive_config.get('table', 'dq_check_results')

query = f"""
SELECT *
FROM {database}.{table}
WHERE rule_id = '{rule_id}'
 AND check_date >= date_sub(current_date(), {days})
ORDER BY check_timestamp DESC
"""

return self.spark.sql(query)

def calculate_trends(self, rule_id: str, days: int = 30) -> Dict:
 """Calculate trends for a specific rule"""
 historical_df = self.get_historical_results(rule_id, days)

 historical_df.createOrReplaceTempView("historical_results")

 trends_query = """
 SELECT
 rule_id,
 check_date,
 AVG(failure_percentage) as avg_failure_rate,
 MIN(failure_percentage) as min_failure_rate,
 MAX(failure_percentage) as max_failure_rate,
 COUNT(*) as check_count
 FROM historical_results
 GROUP BY rule_id, check_date
 ORDER BY check_date
 """
```

```
trends_df = self.spark.sql(trends_query)
return trends_df.toPandas().to_dict('records')
```

## src/notifier/alert\_manager.py

python

```
from typing import Dict, List
import requests
import smtplib
from email.mime.text import MIMEText
from email.mime.multipart import MIMEMultipart
from datetime import datetime, timedelta
import json
import logging

logger = logging.getLogger(__name__)

class AlertManager:
 """Manages alerts across multiple channels"""

 def __init__(self, config: Dict):
 self.config = config
 self.alert_config = config.get('alerting', {})
 self.notification_config = config.get('notification', {})
 self.alert_history = []
 self.last_alert_time = {}

 def send_alert(self, rule_result: Dict, severity: str):
 """Send alert through configured channels based on severity"""
```

```
if not self._should_send_alert(rule_result['rule_id'], severity):
 return

channels = self._get_channels_for_severity(severity)

for channel in channels:
 if channel == 'email' and self.alert_config.get('email', {}).get('enabled'):
 self._send_email_alert(rule_result, severity)
 elif channel == 'slack' and self.alert_config.get('slack', {}).get('enabled'):
 self._send_slack_alert(rule_result, severity)
 elif channel == 'teams' and self.alert_config.get('teams', {}).get('enabled'):
 self._send_teams_alert(rule_result, severity)
 elif channel == 'pagerduty' and self.alert_config.get('pagerduty', {}).get('enabled'):
 self._send_pagerduty_alert(rule_result, severity)

 # Record alert
 self.alert_history.append({
 'timestamp': datetime.now(),
 'rule_id': rule_result['rule_id'],
 'severity': severity,
 'result': rule_result
 })

 # Update last alert time
 self.last_alert_time[rule_result['rule_id']] = datetime.now()

def _should_send_alert(self, rule_id: str, severity: str) -> bool:
 """Check if alert should be sent based on cooldown"""
 if rule_id not in self.last_alert_time:
 return True
```

```
cooldown = self.notification_config.get('cooldown_minutes', 30)
if datetime.now() - self.last_alert_time[rule_id] < timedelta(minutes=cooldown):
 logger.info(f"Alert suppressed for {rule_id} due to cooldown")
 return False

return True

def _get_channels_for_severity(self, severity: str) -> List[str]:
 """Get notification channels for severity level"""
 channels = []

 if self.alert_config.get('email', {}).get('recipients', {}).get(severity):
 channels.append('email')
 if self.alert_config.get('slack', {}).get('channels', {}).get(severity):
 channels.append('slack')
 if severity == 'critical' and self.alert_config.get('pagerduty', {}).get('enable
d'):
 channels.append('pagerduty')

 return channels

def _send_email_alert(self, rule_result: Dict, severity: str):
 """Send email alert"""
 try:
 email_config = self.alert_config['email']
 recipients = email_config['recipients'].get(severity, [])

 if not recipients:
 return

 msg = MIMEMultipart()
 msg['Subject'] = f"[{severity.upper()}] DQ Alert - {rule_result['rule_name']}"
```

```
msg['From'] = email_config.get('from', 'dq-alerts@example.com')
msg['To'] = ', '.join(recipients)

body = self._format_email_body(rule_result, severity)
msg.attach(MIMEText(body, 'html'))

with smtplib.SMTP(email_config['smtp_host'], email_config['smtp_port']) as serv
er:

 if email_config.get('use_tls', True):
 server.starttls()
 if email_config.get('username'):
 server.login(email_config['username'], email_config['password'])
 server.send_message(msg)

logger.info(f"Email alert sent for {rule_result['rule_id']}")

except Exception as e:
 logger.error(f"Failed to send email alert: {e}")

def _send_slack_alert(self, rule_result: Dict, severity: str):
 """Send Slack alert"""
 try:
 slack_config = self.alert_config['slack']
 webhook_url = slack_config['webhook_url']
 channel = slack_config['channels'].get(severity, '#data-alerts')

 message = self._format_slack_message(rule_result, severity)

 payload = {
 'channel': channel,
 'text': message['text'],
 'blocks': message['blocks']
 }
```



```
}

response = requests.post(webhook_url, json=payload)
response.raise_for_status()

logger.info(f"Slack alert sent for {rule_result['rule_id']}")

except Exception as e:
 logger.error(f"Failed to send Slack alert: {e}")

def _send_pagerduty_alert(self, rule_result: Dict, severity: str):
 """Send PagerDuty alert"""
 try:
 pd_config = self.alert_config['pagerduty']

 payload = {
 'routing_key': pd_config['api_key'],
 'event_action': 'trigger',
 'payload': {
 'summary': f"DQ Alert: {rule_result['rule_name']}",
 'severity': severity,
 'source': 'DQ Framework',
 'component': rule_result['rule_id'],
 'custom_details': rule_result
 }
 }

 response = requests.post(
 'https://events.pagerduty.com/v2/enqueue',
 json=payload
)
 response.raise_for_status()
```

```

 logger.info(f"PagerDuty alert sent for {rule_result['rule_id']}")

 except Exception as e:
 logger.error(f"Failed to send PagerDuty alert: {e}")

def _format_email_body(self, rule_result: Dict, severity: str) -> str:
 """Format HTML email body"""
 return f"""
 <html>
 <body>
 <h2 style="color: {'red' if severity == 'critical' else 'orange'};">
 Data Quality Alert - {severity.upper()}
 </h2>
 <h3>{rule_result['rule_name']}</h3>
 <table border="1" cellpadding="5">
 <tr><td>Rule ID</td><td>{rule_result['rule_id']}</td></tr>
 <tr><td>Table</td><td>{rule_result['table_name']}</td></tr>
 <tr><td>Total Rows</td><td>{rule_result['total_rows']}</td></tr>
 <tr><td>Failed Rows</td><td>{rule_result['failed_rows']}</td></tr>
 <tr><td>Failure %</td><td>{rule_result['failure_percentage']}%</td></tr>
 <tr><td>Threshold</td><td>{rule_result.get('threshold', 'N/A')}%</td></tr>
 <tr><td>Timestamp</td><td>{rule_result['check_timestamp']}</td></tr>
 </table>
 <p>Please investigate and take necessary action.</p>
 </body>
 </html>
 """

def _format_slack_message(self, rule_result: Dict, severity: str) -> Dict:
 """Format Slack message with blocks"""
 emoji = {

```

```

 'critical': '🔴',
 'high': '🟠',
 'medium': '🟡',
 'warning': '⚪'
 }

 return {
 'text': f"{emoji.get(severity, '⚪')} DQ Alert: {rule_result['rule_name']}",
 'blocks': [
 {
 'type': 'header',
 'text': {
 'type': 'plain_text',
 'text': f"{emoji.get(severity, '⚪')} Data Quality Alert"
 }
 },
 {
 'type': 'section',
 'fields': [
 {'type': 'mrkdwn', 'text': f"*Rule:* {rule_result['rule_name']}"},
 {'type': 'mrkdwn', 'text': f"*Severity:* {severity.upper()}"},
 {'type': 'mrkdwn', 'text': f"*Table:* {rule_result['table_name']}"},
 {'type': 'mrkdwn', 'text': f"*Failure Rate:* {rule_result['failure_percentage']}%"}
]
 },
 {
 'type': 'context',
 'elements': [
 {
 'type': 'mrkdwn',

```

```

 'text': f"Rule ID: {rule_result['rule_id']} | {rule_result['che
ck_timestamp']}"
 }
]
}

```

## main.py - Main Orchestrator

```

python

#!/usr/bin/env python3
"""
Enterprise Data Quality Framework – Main Orchestrator
"""

import sys
import logging
from pathlib import Path
from datetime import datetime
from typing import Dict, List
import argparse

Add src to path
sys.path.append(str(Path(__file__).parent))

from src.parser.rule_parser import RuleParser
from src.engine.metadata_engine import MetadataEngine
from src.resolver.partition_resolver import PartitionResolver
from src.compiler.rule_compiler import RuleCompiler

```

```
from src.generator.spark_sql_generator import SparkSQLGenerator
from src.executor.spark_executor import SparkExecutor
from src.storage.result_store import ResultStore
from src.notifier.alert_manager import AlertManager

from pyspark.sql import SparkSession
import yaml

logging.basicConfig(
 level=logging.INFO,
 format='%(asctime)s - %(name)s - %(levelname)s - %(message)s',
 handlers=[
 logging.FileHandler('dq_framework.log'),
 logging.StreamHandler(sys.stdout)
]
)

logger = logging.getLogger(__name__)

class DQFramework:
 """Main orchestrator for the Data Quality Framework"""

 def __init__(self, config_path: str):
 self.config_path = Path(config_path)
 self.config = self._load_config()
 self.spark = self._create_spark_session()

 # Initialize components
 self.rule_parser = RuleParser(config_path)
 self.metadata_engine = MetadataEngine(
 self.spark,
 self.config['hive']['metastore_uri']
```

```
)
self.partition_resolver = PartitionResolver(self.config)
self.rule_compiler = RuleCompiler()
self.spark_generator = SparkSQLGenerator(self.spark)
self.spark_executor = SparkExecutor(self.spark, self.config)
self.result_store = ResultStore(self.spark, self.config)
self.alert_manager = AlertManager(self.config)

def _load_config(self) -> Dict:
 """Load main DQ configuration"""
 config_file = self.config_path / 'dq_config.yaml'
 with open(config_file, 'r') as f:
 config = yaml.safe_load(f)
 return config['framework']

def _create_spark_session(self) -> SparkSession:
 """Create and configure Spark session"""
 spark_config = self.config['spark']

 builder = SparkSession.builder \
 .appName(spark_config['app_name']) \
 .master(spark_config.get('master', 'yarn')) \
 .config("spark.sql.adaptive.enabled", "true") \
 .config("spark.sql.adaptive.coalescePartitions.enabled", "true") \
 .config("spark.sql.catalogImplementation", "hive") \
 .enableHiveSupport()

 return builder.getOrCreate()

def run_single_table_check(self, table_yaml: str) -> Dict:
 """Run DQ checks for a single table"""
 logger.info(f"Running DQ checks for table config: {table_yaml}")
```

```
Parse rules
config = self.rule_parser.load_yaml_config(table_yaml)
table_config = config['table']
rules = table_config['data_quality']['rules']

results = []

for rule_config in rules:
 try:
 # Get metadata
 table_meta = self.metadata_engine.get_table_metadata(
 table_config['database'],
 table_config['name']
)

 # Resolve partitions
 partitions = self.partition_resolver.resolve_partitions(
 table_config,
 rule_config
)

 # Generate partition filters
 partition_filters = self.partition_resolver.generate_partition_filters(
 partitions
)

 # For each partition, compile and execute
 for partition_filter in partition_filters:
 # Compile rule
 compiled_rule = self.rule_compiler.compile_rule(
 rule_config,
```

```
 table_config,\n partition_filter\n)\n\n # Generate SQL\n sql_query = self.spark_generator.generate_check_query(compiled_rule)\n\n # Execute\n result_df = self.spark_executor.execute_single_check(sql_query)\n\n # Process results\n result_data = result_df.collect()[0].asDict()\n\n # Evaluate threshold\n failure_pct = result_data.get('failure_percentage', 0)\n threshold = rule_config.get('threshold', 0)\n\n if failure_pct > threshold:\n result_data['status'] = 'FAILED'\n self.alert_manager.send_alert(result_data, rule_config['severity'])\n else:\n result_data['status'] = 'PASSED'\n\n results.append(result_data)\n\nexcept Exception as e:\n logger.error(f"Failed to execute rule {rule_config['rule_id']}: {e}")\n results.append({\n 'rule_id': rule_config['rule_id'],\n 'status': 'ERROR',\n 'error': str(e),\n 'check_timestamp': datetime.now()
```



```
 })

 # Store results
 if results:
 results_df = self.spark.createDataFrame(results)
 self.result_store.store_results(results_df)

 return {'table': table_config['name'], 'results': results}

def run_all_checks(self) -> Dict:
 """Run all DQ checks for all configured tables"""
 logger.info("Starting full DQ check run")

 # Parse all configurations
 rules = self.rule_parser.parse_all_configs()

 all_results = {}

 # Group rules by table for efficient processing
 tables_config = {}
 config_dir = self.config_path / 'tables'
 for yaml_file in config_dir.glob('*.yaml'):
 config = self.rule_parser.load_yaml_config(str(yaml_file))
 table_name = config['table']['name']
 tables_config[table_name] = str(yaml_file)

 # Process each table
 for table_name, config_file in tables_config.items():
 try:
 results = self.run_single_table_check(config_file)
 all_results[table_name] = results
 except Exception as e:
```

```
 logger.error(f"Failed to process table {table_name}: {e}")
 all_results[table_name] = {'status': 'ERROR', 'error': str(e)}

 # Generate summary
 summary = self._generate_summary(all_results)

 # Store summary metrics
 self.result_store.store_metrics(summary['metrics'])

 logger.info(f"DQ check completed. Summary: {summary['overall']}")
 return all_results

def _generate_summary(self, results: Dict) -> Dict:
 """Generate summary of all DQ checks"""
 total_rules = 0
 passed = 0
 failed = 0
 errors = 0
 metrics = []

 for table_name, table_results in results.items():
 for result in table_results.get('results', []):
 total_rules += 1
 status = result.get('status', 'ERROR')

 if status == 'PASSED':
 passed += 1
 elif status == 'FAILED':
 failed += 1
 else:
 errors += 1
```

```
 metrics.append({
 'rule_id': result.get('rule_id'),
 'table_name': table_name,
 'check_timestamp': result.get('check_timestamp'),
 'total_rows': result.get('total_rows', 0),
 'failed_rows': result.get('failed_rows', 0),
 'failure_percentage': result.get('failure_percentage', 0),
 'status': status
 })

 return {
 'overall': {
 'total_rules': total_rules,
 'passed': passed,
 'failed': failed,
 'errors': errors,
 'pass_rate': (passed / total_rules * 100) if total_rules > 0 else 0
 },
 'metrics': metrics
 }

def run_incremental_check(self):
 """Run incremental DQ checks for new data only"""
 logger.info("Running incremental DQ check")
 # Implementation for incremental checks
 pass

def cleanup(self):
 """Cleanup resources"""
 self.spark.stop()

def main():
```

```
parser = argparse.ArgumentParser(description='Data Quality Framework')
parser.add_argument(
 '--config',
 type=str,
 default='config',
 help='Path to configuration directory'
)
parser.add_argument(
 '--table',
 type=str,
 help='Run checks for specific table only'
)
parser.add_argument(
 '--mode',
 type=str,
 choices=['full', 'incremental', 'single'],
 default='full',
 help='Execution mode'
)

args = parser.parse_args()

Initialize framework
dq_framework = DQFramework(args.config)

try:
 if args.mode == 'single' and args.table:
 # Run for single table
 config_file = Path(args.config) / 'tables' / f"{args.table}.yaml"
 if config_file.exists():
 dq_framework.run_single_table_check(str(config_file))
 else:
```

```
 logger.error(f"Config file not found: {config_file}")
 elif args.mode == 'incremental':
 dq_framework.run_incremental_check()
 else:
 # Run full check
 dq_framework.run_all_checks()
 finally:
 dq_framework.cleanup()

if __name__ == "__main__":
 main()
```

### 3. Tests

#### tests/test\_dq\_framework.py

python

```
import unittest
from unittest.mock import Mock, patch
import yaml
from pathlib import Path
import sys

sys.path.append(str(Path(__file__).parent.parent))

from src.parser.rule_parser import RuleParser
from src.compiler.rule_compiler import RuleCompiler
from src.resolver.partition_resolver import PartitionResolver
```

```
class TestRuleParser(unittest.TestCase):
 def setUp(self):
 self.parser = RuleParser('config')

 def test_load_valid_yaml(self):
 yaml_content = """
 table:
 name: test_table
 database: test_db
 data_quality:
 rules:
 - rule_id: TEST_001
 name: Test Rule
 type: null_check
 severity: critical
 columns:
 - test_column
 threshold: 0
 """
 # Mock file reading
 with patch('builtins.open', unittest.mock.mock_open(read_data=yaml_content)):
 config = self.parser.load_yaml_config('test.yaml')
 self.assertEqual(config['table']['name'], 'test_table')

 def test_invalid_yaml_structure(self):
 invalid_yaml = """
 invalid_key: value
 """
 with patch('builtins.open', unittest.mock.mock_open(read_data=invalid_yaml)):
 with self.assertRaises(ValueError):
 self.parser.load_yaml_config('test.yaml')
```

```
class TestRuleCompiler(unittest.TestCase):
 def setUp(self):
 self.compiler = RuleCompiler()
 self.table_config = {
 'database': 'test_db',
 'name': 'test_table'
 }

 def test_compile_null_check(self):
 rule = {
 'rule_id': 'TEST_001',
 'name': 'Null Check',
 'type': 'null_check',
 'severity': 'critical',
 'columns': ['col1', 'col2'],
 'threshold': 0
 }

 compiled = self.compiler.compile_rule(rule, self.table_config)
 self.assertIn('IS NULL', compiled.sql_query)
 self.assertIn('test_db.test_table', compiled.sql_query)

 def test_compile_value_range(self):
 rule = {
 'rule_id': 'TEST_002',
 'name': 'Range Check',
 'type': 'value_range',
 'severity': 'high',
 'column': 'amount',
 'min_value': 0,
 'max_value': 1000,
 'threshold': 0
 }
```

```
}

compiled = self.compiler.compile_rule(rule, self.table_config)
self.assertIn('amount < 0', compiled.sql_query)
self.assertIn('amount > 1000', compiled.sql_query)

class TestPartitionResolver(unittest.TestCase):
 def setUp(self):
 self.config = {
 'partitions': {
 'start_date': '2024-01-01',
 'end_date': '2024-01-05'
 }
 }
 self.resolver = PartitionResolver(self.config)

 def test_resolve_full_partitions(self):
 table_config = {
 'partitions': {
 'type': 'date_hierarchy',
 'columns': ['year', 'month', 'day']
 }
 }

 partitions = self.resolver.resolve_partitions(table_config, {})
 self.assertTrue(len(partitions) > 0)

 def test_generate_partition_filters(self):
 partitions = [
 {'year': '2024', 'month': '01', 'day': '01'},
 {'year': '2024', 'month': '01', 'day': '02'}
]
```



```
 filters = self.resolver.generate_partition_filters(partitions)
 self.assertEqual(len(filters), 2)
 self.assertIn("year='2024'", filters[0])

if __name__ == '__main__':
 unittest.main()
```

## 4. Requirements

### requirements.txt

```
txt

pyspark==3.5.0
pyyaml==6.0.1
jsonschema==4.20.0
psycopg2-binary==2.9.9
pymysql==1.1.0
requests==2.31.0
pandas==2.1.4
dataclasses==0.6
typing-extensions==4.9.0
```

## 5. Usage Examples

### Running the Framework

```
bash

Install dependencies
pip install -r requirements.txt

Run full DQ checks
python main.py --config config --mode full

Run single table check
python main.py --config config --table orders --mode single

Run incremental checks
python main.py --config config --mode incremental

Schedule via cron
0 */6 * * * cd /path/to/dq_framework && python main.py --config config --mode incremental
```

## Docker Deployment

```
dockerfile

Dockerfile
FROM apache/spark-py:v3.5.0

WORKDIR /app

COPY requirements.txt .
RUN pip install -r requirements.txt

COPY . .
```

```
ENTRYPOINT ["python", "main.py"]
CMD ["--config", "config", "--mode", "full"]
```

## Airflow DAG Integration

```
python

airflow_dag.py
from airflow import DAG
from airflow.operators.python_operator import PythonOperator
from datetime import datetime, timedelta

default_args = {
 'owner': 'data-engineering',
 'depends_on_past': False,
 'start_date': datetime(2024, 1, 1),
 'email_on_failure': True,
 'email_on_retry': False,
 'retries': 2,
 'retry_delay': timedelta(minutes=5)
}

dag = DAG(
 'dq_framework',
 default_args=default_args,
 description='Data Quality Checks',
 schedule_interval='0 */6 * * *',
 catchup=False
)

def run_dq_checks(**context):
```

```
from main import DQFramework
dq = DQFramework('config')
dq.run_all_checks()

t1 = PythonOperator(
 task_id='run_dq_checks',
 python_callable=run_dq_checks,
 dag=dag
)
```

## Key Features Implemented:

1. **Modular Architecture:** Each component has a single responsibility
2. **YAML-Based Configuration:** DBT-style configurations with validation
3. **Dynamic Partition Handling:** Supports multiple partition strategies
4. **Rule Types:** Null checks, freshness, completeness, uniqueness, custom SQL, cross-table validation
5. **Scalable Execution:** Parallel execution with Spark optimization
6. **Alerting:** Multi-channel alerts (Email, Slack, Teams, PagerDuty)
7. **Result Storage:** Dual storage in Hive and PostgreSQL
8. **Error Handling:** Comprehensive error handling and retry logic
9. **Monitoring:** Metrics collection and trend analysis
10. **Enterprise Ready:** Production-grade features for large-scale deployments

This framework provides a robust, scalable solution for data quality checks on Hive tables with partition-aware processing and comprehensive alerting capabilities.